


In [33]:

```

1 cm = confusion_matrix(y_test, y_pred_lda)
2
3 df1 = pd.DataFrame(columns=["0", "1", "2", "3"], index= ["0", "1", "2", "3"], data= cm )
4
5 f,ax = plt.subplots(figsize=(6,6))
6
7 sns.heatmap(df1, annot=True,cmap="Greens", fmt= '.0f',ax=ax,linewidths = 5, cbar = False)
8 plt.xlabel("Predicted Label")
9 plt.xticks(size = 12)
10 plt.yticks(size = 12, rotation = 0)
11 plt.ylabel("True Label")
12 plt.title("Confusion Matrix", size = 12)
13 plt.show()

```



## Hyperparameter Tuning - Grid Search - Cross Validation

```

1
2 # We will compare 8 classifier and evaluate mean accuracy of each of them by
  stratified cross validation.
3
4 Decision Tree Classifier
5 SVC
6 Random Forest Classifier
7 Logistic Regression
8 KNN Classifier
9 Stochastic Gradient Descent Classifier
10 Gradient Boosting Classifier
11 LightGBM Classifier

```

In [34]:

```

1 classifier = [DecisionTreeClassifier(random_state = random_state),
2               SVC(random_state = random_state, probability = True),
3               RandomForestClassifier(random_state = random_state),
4               LogisticRegression(random_state = random_state),
5               KNeighborsClassifier(),
6               SGDClassifier(random_state = random_state),
7               GradientBoostingClassifier(random_state = random_state),
8               LGBMClassifier(random_state = random_state)]
9
10 dt_param_grid = {"min_samples_split" : range(10,500,20),
11                  "max_depth": range(1,20,2)}
12
13 svc_param_grid = {"kernel" : ["rbf"],
14                   "gamma": [0.001, 0.01, 0.1, 1],
15                   "C": [1,10,50,100,200,300,1000]}
16
17 rf_param_grid = {"max_features": [1,3,10],
18                  "min_samples_split": [2,3,10],
19                  "min_samples_leaf": [1,3,10],
20                  "bootstrap": [False],
21                  "n_estimators": [100,300],
22                  "criterion": ["gini"]}
23
24 logreg_param_grid = {"C": np.logspace(-4, 4, 20),
25                      "penalty": ["l1", "l2", "none"]}
26
27 knn_param_grid = {"n_neighbors": np.linspace(2,20,12, dtype = int).tolist(),
28                   "weights": ["uniform", "distance"],
29                   "metric": ["euclidean", "manhattan", "minkowski"],
30                   "leaf_size": [30]}
31
32 sgdc_param_grid = {
33     "loss" : ["hinge", "log", "squared_hinge", "modified_huber"],
34     "alpha" : [0.0001, 0.001, 0.01, 0.1],
35     "penalty" : ["l2", "l1", "none"]}
36
37 gbc_param_grid = {
38     "learning_rate": [0.05, 0.1, 0.2],
39     "min_samples_split": [2,3,10],
40     "min_samples_leaf": [1,3,10]
41 }
42
43
44 lgbmc_param_grid = {
45     'num_leaves': [31, 127],
46     'reg_alpha': [0.1, 0.5]}
47
48
49 classifier_param = [dt_param_grid,
50                     svc_param_grid,
51                     rf_param_grid,
52                     logreg_param_grid,
53                     knn_param_grid,
54                     sgdc_param_grid,
55                     gbc_param_grid,
56                     lgbmc_param_grid]
57
58 cv_result = []
59 best_estimators = []

```

```

60 mean_squared_errors = []
61 roc_auc_scores = []
62 recall_scores = []
63 precision_scores = []
64 f1_scores = []
65
66
67 for i in range(len(classifier)):
68     print("-----")
69     clf = GridSearchCV(classifier[i],
70                        param_grid=classifier_param[i],
71                        cv = StratifiedKFold(n_splits = 10),
72                        scoring = "accuracy",
73                        n_jobs = -1, verbose = 2)
74
75     clf.fit(X_train,y_train)
76
77     cv_result.append(clf.best_score_)
78
79     mean_squared_errors.append(mean_squared_error(y_test,clf.predict(X_test)))
80
81     roc_auc_scores.append(roc_auc_score(y_test, clf.predict_proba(X_test), multi_class
82
83     recall_scores.append(recall_score(y_test, clf.predict(X_test), average='weighted')
84
85     precision_scores.append(precision_score(y_test, clf.predict(X_test), average='weig
86
87     f1_scores.append(f1_score(y_test, clf.predict(X_test), average='weighted'))
88
89     best_estimators.append(clf.best_estimator_)
90
91     print("Model: {}".format(classifier[i]))
92     print("Accuracy: {}".format(round(cv_result[i]*100,2)))
93     print("MSE: {}".format(mean_squared_errors[i]))
94     print("ROC AUC: {}".format(roc_auc_scores[i]))
95     print("Recall: {}".format(recall_scores[i]))
96     print("Precision: {}".format(precision_scores[i]))
97     print("F1-Score: {}".format(f1_scores[i]))
98     print("Best Estimator: {}".format(clf.best_estimator_))
99
100 print("-----")
101
102 sns.set_style("darkgrid")
103 cv_results = pd.DataFrame({"Accuracy":cv_result,
104                           "MSE":mean_squared_errors,
105                           "ROC AUC":roc_auc_scores,
106                           "Recall": recall_scores,
107                           "Precision": precision_scores,
108                           "F1-Score":f1_scores,
109                           "Models":["DecisionTreeClassifier",
110                                    "SVC",
111                                    "RandomForestClassifier",
112                                    "LogisticRegression",
113                                    "KNeighborsClassifier",
114                                    "SGDClassifier",
115                                    "GBClassifier",
116                                    "LGBMClassifier"]})
117
118 cv_results.index = cv_results["Models"]
119
120 cv_results = cv_results.drop(["Models"], axis = 1)

```

```
121
122 f,ax = plt.subplots(figsize=(14,10))
123
124 sns.heatmap(cv_results, annot=True,cmap = "Blues",fmt= '.3f',
125             ax=ax,linewidths = 5, cbar = False,
126             annot_kws={"size": 18})
127
128 plt.xticks(size = 18)
129 plt.yticks(size = 18, rotation = 0)
130 plt.ylabel("Models")
131 plt.title("Grid Search Results", size = 16)
132 plt.show()
```

-----  
-  
Fitting 10 folds for each of 250 candidates, totalling 2500 fits

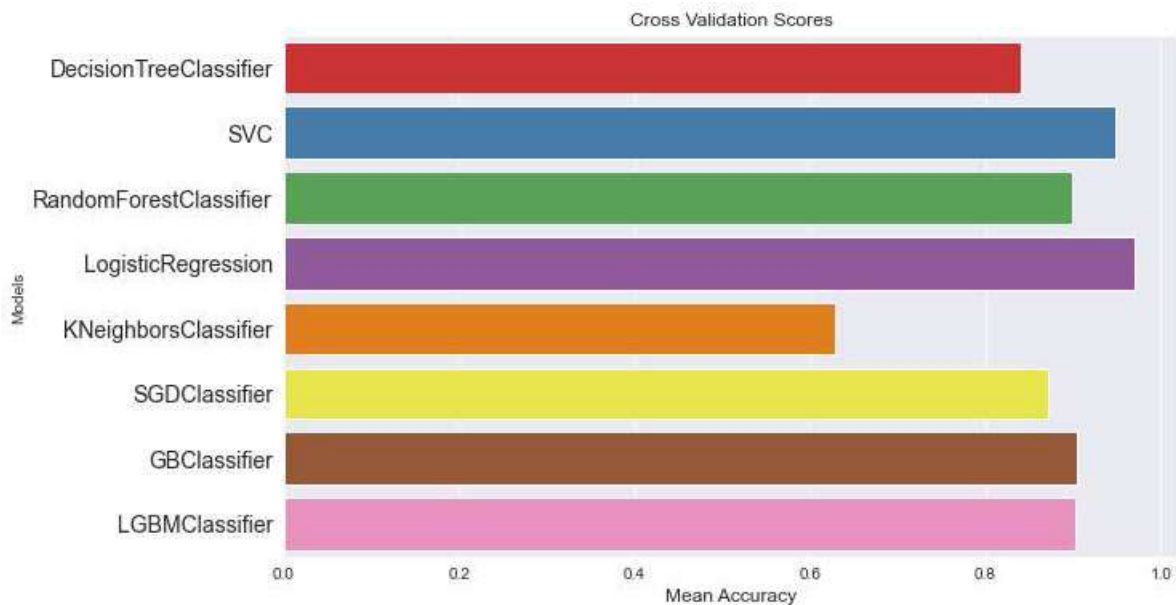
```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent worker
s.
[Parallel(n_jobs=-1)]: Done 72 tasks      | elapsed:    0.2s
[Parallel(n_jobs=-1)]: Done 2408 tasks   | elapsed:    5.4s
[Parallel(n_jobs=-1)]: Done 2500 out of 2500 | elapsed:    5.4s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent worker
s.
```

```
Model: DecisionTreeClassifier(random_state=42)
Accuracy: %84.06
MSE: 0.175
ROC AUC: 0.911858935204842
Recall: 0.825
Precision: 0.8276565656565658
F1-Score: 0.8257884190113292
```

## Cross Validation Scores

In [35]:

```
1 sns.set_style("darkgrid")
2 cv_results = pd.DataFrame({"Cross Validation Means":cv_result,
3                             "Models":["DecisionTreeClassifier", "SVC",
4                                         "RandomForestClassifier",
5                                         "LogisticRegression",
6                                         "KNeighborsClassifier",
7                                         "SGDClassifier",
8                                         "GBClassifier",
9                                         "LGBMClassifier"]})
10
11 plt.figure(figsize = (10,6))
12 sns.barplot("Cross Validation Means", "Models",
13             data = cv_results, palette = "Set1")
14 plt.xlabel("Mean Accuracy",
15             size = 12)
16 plt.yticks(size = 14)
17 plt.title("Cross Validation Scores",
18           size = 12)
19 plt.show()
```



## Ensemble Learning