

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

In [2]:

```
1 import nltk
2 from nltk.sentiment.vader import SentimentIntensityAnalyzer
3
4 nltk.download('vader_lexicon')
5 sid = SentimentIntensityAnalyzer()
```

```
[nltk_data] Downloading package vader_lexicon to
[nltk_data] C:\Users\MSI\AppData\Roaming\nltk_data...
[nltk_data] Package vader_lexicon is already up-to-date!
```

VADER's `SentimentIntensityAnalyzer()` takes in a string and returns a dictionary of scores in each of four categories:

- Negative
- Neutral
- Positive Compound (computed by normalizing the scores above)

In [3]:

```
1 a = 'This was a good movie'
2 b = 'The was the best, most awesome movie EVER MADE!!!'
3 c = 'This was the worst film to ever disgrace the screen.'
4
5 print(sid.polarity_scores(a))
6 print(sid.polarity_scores(b))
7 print(sid.polarity_scores(c))
```

```
{'neg': 0.0, 'neu': 0.508, 'pos': 0.492, 'compound': 0.4404}
{'neg': 0.0, 'neu': 0.425, 'pos': 0.575, 'compound': 0.8877}
{'neg': 0.477, 'neu': 0.523, 'pos': 0.0, 'compound': -0.8074}
```

Use VADER to analyze Amazon Reviews

In [4]:

```
1 data = pd.read_csv('IMDB Dataset.csv')
2 data.head()
```

Out[4]:


	review	sentiment
0	One of the other reviewers has mentioned that ...	positive
1	A wonderful little production. The...	positive
2	I thought this was a wonderful way to spend ti...	positive
3	Basically there's a family where a little boy ...	negative
4	Petter Mattei's "Love in the Time of Money" is...	positive

In [5]:

```
1 sns.countplot(x='sentiment', data=data)
```

Out[5]:

<AxesSubplot:xlabel='sentiment', ylabel='count'>



In [6]:

```
1 data['scores'] = data['review'].apply(lambda review: sid.polarity_scores(review))
```

In [7]:

```
1 data.head()
```

Out[7]:

	review	sentiment	scores
0	One of the other reviewers has mentioned that ...	positive	{'neg': 0.203, 'neu': 0.748, 'pos': 0.048, 'co...
1	A wonderful little production. The...	positive	{'neg': 0.053, 'neu': 0.776, 'pos': 0.172, 'co...
2	I thought this was a wonderful way to spend ti...	positive	{'neg': 0.094, 'neu': 0.714, 'pos': 0.192, 'co...
3	Basically there's a family where a little boy ...	negative	{'neg': 0.138, 'neu': 0.797, 'pos': 0.065, 'co...
4	Petter Mattei's "Love in the Time of Money" is...	positive	{'neg': 0.052, 'neu': 0.801, 'pos': 0.147, 'co...

In [8]:

```
1 data['compound'] = data['scores'].apply(lambda score_dict: score_dict['compound'])
2 data['comp_score'] = data['compound'].apply(lambda c: 'positive' if c >= 0 else 'negative')
```

In [9]:

```
1 data.head()
```

Out[9]:

	review	sentiment	scores	compound	comp_score
0	One of the other reviewers has mentioned that ...	positive	{'neg': 0.203, 'neu': 0.748, 'pos': 0.048, 'co...	-0.9951	negative
1	A wonderful little production. The...	positive	{'neg': 0.053, 'neu': 0.776, 'pos': 0.172, 'co...	0.9641	positive
2	I thought this was a wonderful way to spend ti...	positive	{'neg': 0.094, 'neu': 0.714, 'pos': 0.192, 'co...	0.9605	positive
3	Basically there's a family where a little boy ...	negative	{'neg': 0.138, 'neu': 0.797, 'pos': 0.065, 'co...	-0.9213	negative
4	Petter Mattei's "Love in the Time of Money" is...	positive	{'neg': 0.052, 'neu': 0.801, 'pos': 0.147, 'co...	0.9744	positive

In [10]:

```
1 from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
2
3 accuracy_score(data['sentiment'], data['comp_score'])
```

Out[10]:

0.69556

In [11]:

```
1 confusion_matrix(data['sentiment'], data['comp_score'])
```

Out[11]:

```
array([[13364, 11636],
       [ 3586, 21414]], dtype=int64)
```

In [12]:

```
1 pd.DataFrame(classification_report(data['sentiment'], data['comp_score'], output_dict=True))
```

Out[12]:

	negative	positive	accuracy	macro avg	weighted avg
precision	0.788437	0.647927	0.69556	0.718182	0.718182
recall	0.534560	0.856560	0.69556	0.695560	0.695560
f1-score	0.637139	0.737778	0.69556	0.687459	0.687459
support	25000.000000	25000.000000	0.69556	50000.000000	50000.000000

In [13]:

```
1 data = pd.read_csv('train.csv')
```

In [14]:

```
1 print(data.target)
```

```
0      1
1      1
2      1
3      1
4      1
..
7608   1
7609   1
7610   1
7611   1
7612   1
```

```
Name: target, Length: 7613, dtype: int64
```

Classic Machine Learning Models

In [15]:

```
1 import spacy
2 import pandas as pd
3 import numpy as np
4
5 from sklearn.model_selection import train_test_split
6 from sklearn.feature_extraction.text import TfidfVectorizer
7
8 from sklearn.linear_model import LogisticRegression
9 from sklearn.svm import LinearSVC
10 from sklearn.naive_bayes import MultinomialNB
11
12 from sklearn.metrics import accuracy_score, classification_report
13
14 def train_model(model_name, model, X_train, X_test, y_train, y_test):
15     print(f'BEGIN. {model_name.upper()}.....')
16     model.fit(X_train, y_train)
17     y_pred = model.predict(X_test)
18     y_train_pred = model.predict(X_train)
19     print(f'TESTING DATA----> {model_name.upper()}: \t\t{accuracy_score(y_test, y_pred)}')
20     print(f'TRAINING DATA---> {model_name.upper()}: \t\t{accuracy_score(y_train, y_train_pred)}')
21     print(classification_report(y_test, y_pred))
22     print(f'END. {model_name.upper()}')
23     print('=====')
24     return y_pred
25
26 data = pd.read_csv('train.csv')
```

In [16]:

```

1 print('====Splitting the data====')
2 X = data.text
3 y = data.target
4 print(f'Data shape: {data.shape}')
5
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
7 print(f'X_Train shape: {X_train.shape}, y_train shape: {y_train.shape}')
8 print(f'X_Test shape: {X_test.shape}, y_test shape: {y_test.shape}')
9
10 print('\n====Message Preprocessing====')
11 vectorizer = TfidfVectorizer(ngram_range=(1, 2), max_df=0.9, min_df=2, stop_words='english')
12 X_train_vect = vectorizer.fit_transform(X_train.astype('U'))
13 X_test_vect = vectorizer.transform(X_test)
14
15 print('Training and testing data shape after pre-processing:')
16 print(f'X_Train shape: {X_train_vect.shape}, y_train shape: {y_train.shape}')
17 print(f'X_Test shape: {X_test_vect.shape}, y_test shape: {y_test.shape}')
18
19 print('\n====Model Building====')
20 lr_model = LogisticRegression()
21 lr_y_pred = train_model('Logistic Regression', lr_model, X_train_vect, X_test_vect, y_train, y_test)
22
23 svm_model = LinearSVC()
24 svm_y_pred = train_model('Support Vector Machine', svm_model, X_train_vect, X_test_vect, y_train, y_test)
25
26 nb_model = MultinomialNB()
27 nb_y_pred = train_model('Naive Bayes', nb_model, X_train_vect, X_test_vect, y_train, y_test)
28

```

====Splitting the data====

Data shape: (7613, 5)

X_Train shape: (5329,), y_train shape: (5329,)

X_Test shape: (2284,), y_test shape: (2284,)

====Message Preprocessing====

Training and testing data shape after pre-processing:

X_Train shape: (5329, 8557), y_train shape: (5329,)

X_Test shape: (2284, 8557), y_test shape: (2284,)

====Model Building====

BEGIN. LOGISTIC REGRESSION.....

TESTING DATA----> LOGISTIC REGRESSION: 80.12%

TRAINING DATA---> LOGISTIC REGRESSION: 88.40%

	precision	recall	f1-score	support
0	0.80	0.88	0.84	1318
1	0.81	0.69	0.75	966
accuracy			0.80	2284
macro avg	0.80	0.79	0.79	2284
weighted avg	0.80	0.80	0.80	2284

END. LOGISTIC REGRESSION

====

BEGIN. SUPPORT VECTOR MACHINE.....

TESTING DATA----> SUPPORT VECTOR MACHINE: 77.76%

TRAINING DATA---> SUPPORT VECTOR MACHINE: 95.91%

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.80	0.82	0.81	1318
1	0.75	0.72	0.73	966
accuracy			0.78	2284
macro avg	0.77	0.77	0.77	2284
weighted avg	0.78	0.78	0.78	2284

END. SUPPORT VECTOR MACHINE

=====

BEGIN. NAIVE BAYES.....

TESTING DATA----> NAIVE BAYES: 80.47%

TRAINING DATA---> NAIVE BAYES: 87.15%

	precision	recall	f1-score	support
0	0.78	0.91	0.84	1318
1	0.85	0.66	0.74	966
accuracy			0.80	2284
macro avg	0.82	0.79	0.79	2284
weighted avg	0.81	0.80	0.80	2284

END. NAIVE BAYES

=====

Recurrent Neural Networks

In [17]:

```

1 import pandas as pd
2 import numpy as np
3
4 from sklearn.model_selection import train_test_split
5 from sklearn.feature_extraction.text import TfidfVectorizer
6
7 from sklearn.linear_model import LogisticRegression
8 from sklearn.svm import LinearSVC
9 from sklearn.naive_bayes import MultinomialNB
10
11 from sklearn.metrics import accuracy_score, classification_report
12
13 def train_model(model_name, model, X_train, X_test, y_train, y_test):
14     print(f'BEGIN. {model_name.upper()}.....')
15     model.fit(X_train, y_train)
16     y_pred = model.predict(X_test)
17     y_train_pred = model.predict(X_train)
18     print(f'TESTING DATA----> {model_name.upper()}: \t\t{accuracy_score(y_test, y_pred)}')
19     print(f'TRAINING DATA---> {model_name.upper()}: \t\t{accuracy_score(y_train, y_train_pred)}')
20     print(classification_report(y_test, y_train_pred))
21     print(f'END. {model_name.upper()}')
22     print('=====')
23     return y_pred
24
25 data = pd.read_csv('train.csv')
26
27 print('=====Splitting the data=====')
28 X = data.text
29 y = data.target
30 print(f'Data shape: {data.shape}')
31
32 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
33 print(f'X_Train shape: {X_train.shape}, y_train shape: {y_train.shape}')
34 print(f'X_Test shape: {X_test.shape}, y_test shape: {y_test.shape}')
35
36 print('\n=====Message Preprocessing=====')
37 vectorizer = TfidfVectorizer(ngram_range=(1, 2), max_df=0.9, min_df=2, stop_words='english')
38 X_train_vect = vectorizer.fit_transform(X_train)
39 X_test_vect = vectorizer.transform(X_test)
40
41 print('Training and testing data shape after pre-processing:')
42 print(f'X_Train shape: {X_train_vect.shape}, y_train shape: {y_train.shape}')
43 print(f'X_Test shape: {X_test_vect.shape}, y_test shape: {y_test.shape}')
44
45 print('\n=====Model Building=====')
46 lr_model = LogisticRegression()
47 lr_y_pred = train_model('Logistic Regression', lr_model, X_train_vect, X_test_vect, y_train, y_test)
48
49 svm_model = LinearSVC()
50 svm_y_pred = train_model('Support Vector Machine', svm_model, X_train_vect, X_test_vect, y_train, y_test)
51
52 nb_model = MultinomialNB()
53 nb_y_pred = train_model('Naive Bayes', nb_model, X_train_vect, X_test_vect, y_train, y_test)

```

```

=====Splitting the data=====
Data shape: (7613, 5)
X_Train shape: (5329,), y_train shape: (5329,)
X_Test shape: (2284,), y_test shape: (2284,)

```


=====Message Preprocessing=====

Training and testing data shape after pre-processing:

X_Train shape: (5329, 8557), y_train shape: (5329,)

X_Test shape: (2284, 8557), y_test shape: (2284,)

=====Model Building=====

BEGIN. LOGISTIC REGRESSION.....

TESTING DATA----> LOGISTIC REGRESSION: 80.12%

TRAINING DATA---> LOGISTIC REGRESSION: 88.40%

	precision	recall	f1-score	support
0	0.80	0.88	0.84	1318
1	0.81	0.69	0.75	966
accuracy			0.80	2284
macro avg	0.80	0.79	0.79	2284
weighted avg	0.80	0.80	0.80	2284

END. LOGISTIC REGRESSION

=====

BEGIN. SUPPORT VECTOR MACHINE.....

TESTING DATA----> SUPPORT VECTOR MACHINE: 77.76%

TRAINING DATA---> SUPPORT VECTOR MACHINE: 95.91%

	precision	recall	f1-score	support
0	0.80	0.82	0.81	1318
1	0.75	0.72	0.73	966
accuracy			0.78	2284
macro avg	0.77	0.77	0.77	2284
weighted avg	0.78	0.78	0.78	2284

END. SUPPORT VECTOR MACHINE

=====

BEGIN. NAIVE BAYES.....

TESTING DATA----> NAIVE BAYES: 80.47%

TRAINING DATA---> NAIVE BAYES: 87.15%

	precision	recall	f1-score	support
0	0.78	0.91	0.84	1318
1	0.85	0.66	0.74	966
accuracy			0.80	2284
macro avg	0.82	0.79	0.79	2284
weighted avg	0.81	0.80	0.80	2284

END. NAIVE BAYES

=====

Recurrent Neural Networks

In [18]:

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.model_selection import train_test_split
5
6 from tensorflow.keras.preprocessing.text import Tokenizer
7 from tensorflow.keras.preprocessing.sequence import pad_sequences
8 from tensorflow.keras.layers import Dense, Input, GlobalAveragePooling1D, Dropout, Spa
9 from tensorflow.keras.layers import LSTM, Embedding
10 from tensorflow.keras.models import Model
11 from tensorflow.keras.optimizers import Adam
12
13 from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
14
15
16 def plot_loss_evaluation(r):
17     plt.figure(figsize=(12, 8))
18
19     plt.subplot(2, 2, 1)
20     plt.plot(r.history['loss'], label='loss')
21     plt.plot(r.history['val_loss'], label='val_loss')
22     plt.legend()
23
24     plt.subplot(2, 2, 2)
25     plt.plot(r.history['accuracy'], label='accuracy')
26     plt.plot(r.history['val_accuracy'], label='val_acc')
27     plt.legend()
28
29     plt.title('Training and Loss fuction evolution')
30
31 def evaluate(model, X_train, X_test, y_train, y_test):
32     y_pred_train = np.round(model.predict(X_train))
33     y_pred_test = np.round(model.predict(X_test))
34
35     print("====Training Data====")
36     print(confusion_matrix(y_train, y_pred_train))
37     print(classification_report(y_train, y_pred_train))
38     print(f"Accuracy score: {accuracy_score(y_train, y_pred_train) * 100:.2f}%")
39
40     print("====Testing Data====")
41     print(confusion_matrix(y_test, y_pred_test))
42     print(classification_report(y_test, y_pred_test))
43     print(f"Accuracy score: {accuracy_score(y_test, y_pred_test) * 100:.2f}%")
44
45 data = pd.read_csv("train.csv")
46
47 print('====Splitting the data====')
48 X = data.text
49 y = data.target
50 print(f'Data shape: {data.shape}')
51
52 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=
53 print(f'X_Train shape: {X_train.shape}, y_train shape: {y_train.shape}')
54 print(f'X_Test shape: {X_test.shape}, y_test shape: {y_test.shape}')
55
56 print('====Convert Sentences to Sequences====')
57 MAX_VOCAB_SIZE = 20000
58 tokenizer = Tokenizer(num_words=MAX_VOCAB_SIZE, char_level=False)
59 tokenizer.fit_on_texts(X_train)

```

```

60 sequences_train = tokenizer.texts_to_sequences(X_train)
61 sequences_test = tokenizer.texts_to_sequences(X_test)
62
63 # pad sequence do that we get a NxT matrix
64 data_train = pad_sequences(sequences_train)
65 data_test = pad_sequences(sequences_test, maxlen=data_train.shape[1])
66 print(f"Found {len(tokenizer.word_index)} unique tokens.")
67 print(f"Training Data shape: {data_train.shape}")
68 print(f"Testing Data shape: {data_test.shape}")
69
70 print('=====Create The Model=====')
71 # We get to choose embedding dimensionality
72 D = 100
73 # Hidden state dimentionalitiy
74 M = 64
75 V = len(tokenizer.word_index)
76 T = data_train.shape[1]
77
78 # model.add(embedding)
79 # model.add(SpatialDropout1D(0.2))
80 # model.add(LSTM(64, dropout=0.2, recurrent_dropout=0.2))
81 # model.add(Dense(1, activation='sigmoid'))
82
83 i = Input(shape=(T,))
84 x = Embedding(V + 1, D)(i)
85 x = SpatialDropout1D(0.2)(x)
86 x = LSTM(M, return_sequences=True, activation='relu')(x)
87 x = GlobalAveragePooling1D()(x)
88 # x = Dropout(0.2)(x)
89 x = Dense(1, activation='sigmoid')(x)
90
91 model = Model(i, x)
92 optimizer = Adam(learning_rate=1e-5)
93 # Compile and fit
94 model.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['accuracy'])
95 print('Training model.....')
96 r = model.fit(data_train, y_train, epochs=50,
97               validation_data=(data_test, y_test),
98               batch_size=16)
99
100 print('=====Model Evaluation=====')
101 evaluate(model, data_train, data_test, y_train, y_test)
102 plot_loss_evaluation(r)
103

```

```

=====Splitting the data=====
Data shape: (7613, 5)
X_Train shape: (5329,), y_train shape: (5329,)
X_Test shape: (2284,), y_test shape: (2284,)
=====Convert Sentences to Sequences=====
Found 17762 unique tokens.
Training Data shape: (5329, 33)
Testing Data shape: (2284, 33)
=====Create The Model=====
Training model.....
Epoch 1/50
334/334 [=====] - 18s 44ms/step - loss: 0.6939 -
accuracy: 0.4495 - val_loss: 0.6920 - val_accuracy: 0.5771
Epoch 2/50
334/334 [=====] - 10s 29ms/step - loss: 0.6917 -

```

```
accuracy: 0.5701 - val_loss: 0.6896 - val_accuracy: 0.5771  
Epoch 3/50  
334/334 [=====] - 10s 31ms/step - loss: 0.6899 -  
accuracy: 0.5596 - val loss: 0.6869 - val accuracy: 0.5771
```

In [19]:

```
1 data.text.str.len()
```

Out[19]:

```
0      69  
1      38  
2     133  
3      65  
4      88  
...  
7608    83  
7609   125  
7610    65  
7611   137  
7612    94
```

Name: text, Length: 7613, dtype: int64

Convolutional Neural Networks

In [20]:

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.model_selection import train_test_split
5
6 from tensorflow.keras.preprocessing.text import Tokenizer
7 from tensorflow.keras.preprocessing.sequence import pad_sequences
8 from tensorflow.keras.layers import Dense, Input, GlobalMaxPooling1D, MaxPooling1D
9 from tensorflow.keras.layers import Conv1D, Embedding, Dropout
10 from tensorflow.keras.models import Model
11
12 from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
13
14
15 def plot_loss_evaluation(r):
16     plt.figure(figsize=(12, 8))
17
18     plt.subplot(2, 2, 1)
19     plt.plot(r.history['loss'], label='loss')
20     plt.plot(r.history['val_loss'], label='val_loss')
21     plt.legend()
22
23     plt.subplot(2, 2, 2)
24     plt.plot(r.history['accuracy'], label='accuracy')
25     plt.plot(r.history['val_accuracy'], label='val_acc')
26     plt.legend()
27
28     plt.title('Training and Loss fuction evolution')
29
30 def evaluate(model, X_train, X_test, y_train, y_test):
31     y_pred_train = np.round(model.predict(X_train))
32     y_pred_test = np.round(model.predict(X_test))
33
34     print("====Training Data====")
35     print(confusion_matrix(y_train, y_pred_train))
36     print(classification_report(y_train, y_pred_train))
37     print(f"Accuracy score: {accuracy_score(y_train, y_pred_train) * 100:.2f}%")
38
39     print("====Testing Data====")
40     print(confusion_matrix(y_test, y_pred_test))
41     print(classification_report(y_test, y_pred_test))
42     print(f"Accuracy score: {accuracy_score(y_test, y_pred_test) * 100:.2f}%")
43
44 data = pd.read_csv("train.csv")
45
46 print('====Splitting the data====')
47 X = data.text
48 y = data.target
49 print(f'Data shape: {data.shape}')
50
51 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=
52 print(f'X_Train shape: {X_train.shape}, y_train shape: {y_train.shape}')
53 print(f'X_Test shape: {X_test.shape}, y_test shape: {y_test.shape}')
54
55 print('====Convert Sentences to Sequences====')
56 MAX_VOCAB_SIZE = 20000
57 tokenizer = Tokenizer(num_words=MAX_VOCAB_SIZE)
58 tokenizer.fit_on_texts(X_train)
59 sequences_train = tokenizer.texts_to_sequences(X_train)

```

```

60 sequences_test = tokenizer.texts_to_sequences(X_test)
61
62 # pad sequence do that we get a NxT matrix
63 data_train = pad_sequences(sequences_train)
64 data_test = pad_sequences(sequences_test, maxlen=data_train.shape[1])
65 print(f"Found {len(tokenizer.word_index)} unique tokens.")
66 print(f"Training Data shape: {data_train.shape}")
67 print(f"Testing Data shape: {data_test.shape}")
68
69 print('=====Create The Model=====')
70 # We get to choose embedding dimensionality
71 D = 100
72
73 V = len(tokenizer.word_index)
74 T = data_train.shape[1]
75
76 i = Input(shape=(T,))
77 x = Embedding(V + 1, D)(i)
78
79 x = Conv1D(32, 2, activation='relu')(x)
80 x = MaxPooling1D()(x)
81 x = Dropout(0.1)(x)
82
83 x = Conv1D(64, 2, activation='relu')(x)
84 x = MaxPooling1D()(x)
85 x = Dropout(0.2)(x)
86
87 x = Conv1D(128, 2, activation='relu')(x)
88 x = MaxPooling1D()(x)
89 x = Dropout(0.3)(x)
90
91 x = Conv1D(264, 2, activation='relu')(x)
92 x = GlobalMaxPooling1D()(x)
93
94 x = Dropout(0.5)(x)
95
96 x = Dense(1, activation='sigmoid')(x)
97
98 model = Model(i, x)
99
100 # Compile and fit
101 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
102 print('Training model.....')
103 r = model.fit(data_train, y_train, epochs=5,
104               validation_data=(data_test, y_test),
105               batch_size=1)
106
107 print('=====Model Evaluation=====')
108 evaluate(model, data_train, data_test, y_train, y_test)
109 plot_loss_evaluation(r)

```

```

=====Splitting the data=====
Data shape: (7613, 5)
X_Train shape: (5329,), y_train shape: (5329,)
X_Test shape: (2284,), y_test shape: (2284,)
=====Convert Sentences to Sequences=====
Found 17762 unique tokens.
Training Data shape: (5329, 33)
Testing Data shape: (2284, 33)
=====Create The Model=====

```

Training model.....

Epoch 1/5

5329/5329 [=====] - 134s 25ms/step - loss: 0.6245 - accuracy: 0.6507 - val_loss: 0.4736 - val_accuracy: 0.7973

Epoch 2/5

5329/5329 [=====] - 133s 25ms/step - loss: 0.3623 - accuracy: 0.8628 - val_loss: 0.4646 - val_accuracy: 0.8052

Epoch 3/5

5329/5329 [=====] - 124s 23ms/step - loss: 0.2426 - accuracy: 0.9193 - val_loss: 0.5639 - val_accuracy: 0.7706

Epoch 4/5

5329/5329 [=====] - 131s 25ms/step - loss: 0.1615 - accuracy: 0.9457 - val_loss: 0.6227 - val_accuracy: 0.7842

Epoch 5/5

5329/5329 [=====] - 127s 24ms/step - loss: 0.1178 - accuracy: 0.9615 - val_loss: 0.9541 - val_accuracy: 0.7469

=====Model Evaluation=====

=====Training Data=====

[[2972 52]

[55 2250]]

	precision	recall	f1-score	support
0	0.98	0.98	0.98	3024
1	0.98	0.98	0.98	2305
accuracy			0.98	5329
macro avg	0.98	0.98	0.98	5329
weighted avg	0.98	0.98	0.98	5329

Accuracy score: 97.99%


=====Testing Data=====

[[981 337]

[241 725]]

	precision	recall	f1-score	support
0	0.80	0.74	0.77	1318
1	0.68	0.75	0.71	966
accuracy			0.75	2284
macro avg	0.74	0.75	0.74	2284
weighted avg	0.75	0.75	0.75	2284

Accuracy score: 74.69%



NLTK Sentiment VADER

In [21]:

```

1 from nltk.sentiment.vader import SentimentIntensityAnalyzer
2 sid = SentimentIntensityAnalyzer()
3 data['scores'] = data['text'].apply(lambda text: sid.polarity_scores(text))
4 data['compound'] = data['scores'].apply(lambda score_dict: score_dict['compound'])
5 data['comp_score'] = data['compound'].apply(lambda c: 0 if c >= 0 else 1)

```

In [22]:

```
1 accuracy_score(data['target'], data['comp_score'])
```

Out[22]:

0.5724418757388677

Making submission

In [23]:

```

1 # /kaggle/input/nlp-getting-started/test.csv
2 # /kaggle/input/nlp-getting-started/sample_submission.csv
3 test = pd.read_csv('test.csv')
4
5 print('====Convert Sentences to Sequences====')
6 sequences_test = tokenizer.texts_to_sequences(test.text)
7
8 # pad sequence do that we get a NxT matrix
9 data_test = pad_sequences(sequences_test, maxlen=data_train.shape[1])
10 print(f"Found {len(tokenizer.word_index)} unique tokens.")
11 print(f"Testing Data shape: {data_test.shape}")

```

```
====Convert Sentences to Sequences====
```

```
Found 17762 unique tokens.
```

```
Testing Data shape: (3263, 33)
```

In [24]:

```

1 sample_sub=pd.read_csv('sample_submission.csv')
2 y_pre = model.predict(data_test)
3 y_pre = np.round(y_pre).astype(int).reshape(3263)
4 sub = pd.DataFrame({'id':sample_sub['id'].values.tolist(), 'target':y_pre})
5 sub.to_csv('submission.csv', index=False)

```


In [25]:

```
1 sub.head()
```

Out[25]:

	id	target
0	0	0
1	2	1
2	3	1
3	9	0
4	11	1

In []:

```
1
```